One local residual read — operator detail

Proposed · 768 channels · 8 heads · one output per child

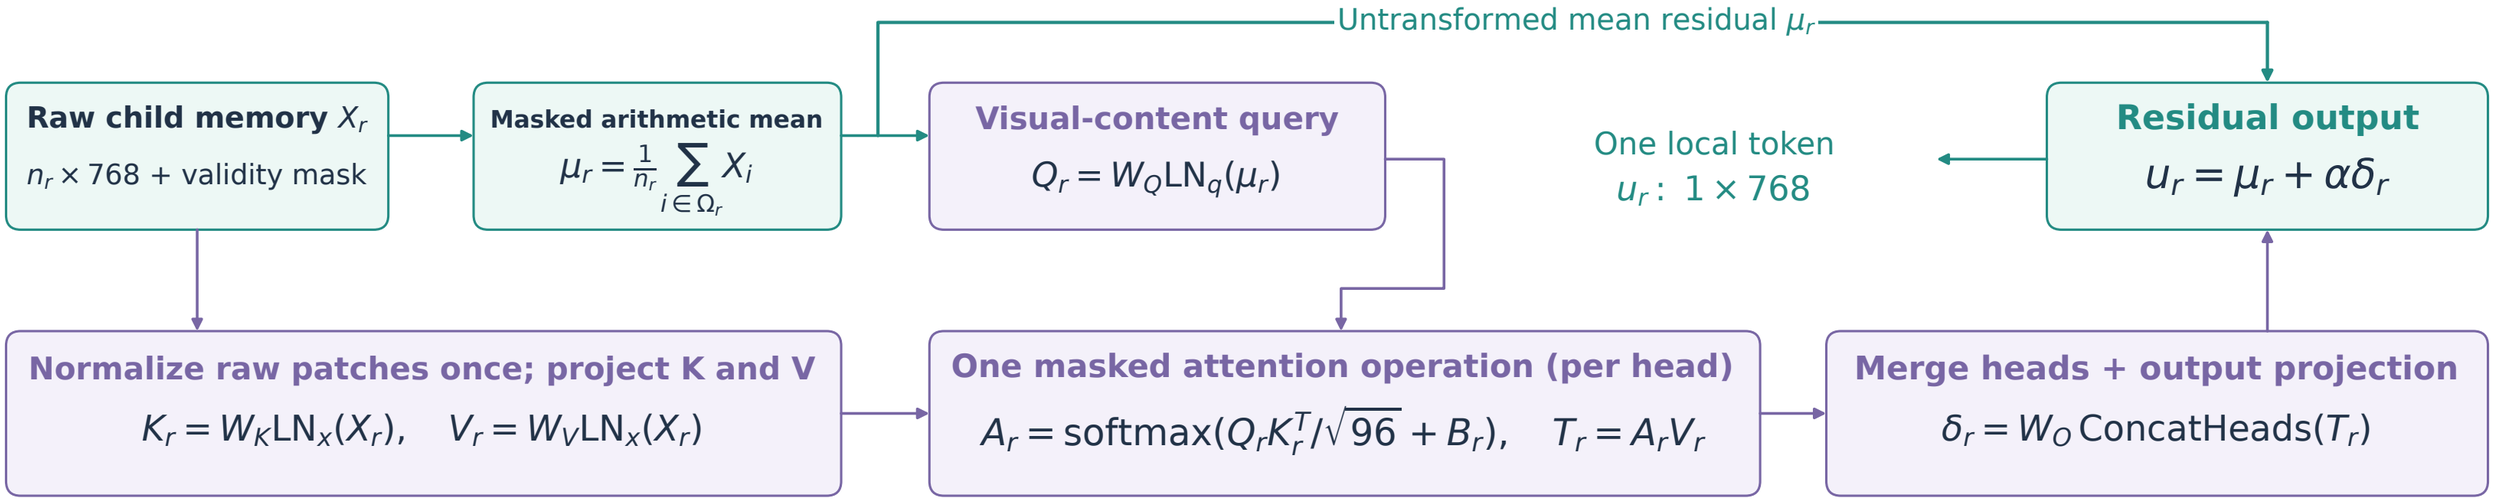

What is preserved

Raw mean bypass · same child identity / order · same total token count

Global branch + final shared LN / MLP · existing report CE only

α: one shared learned scalar, initialized to 0.1 (a proposed starting value)

What this read does not provide by itself

No cross-child raw memory · no explicit relative-position bias

No anatomical labels · no multi-layer encoder input · no local FFN

Attention can still become uniform or select irrelevant content.

$B_r = 0$ for valid child patches; $-\infty$ for packed padding. Shared reader weights; no lesion mask.